

```
import pandas as pd
import numpy as np
from sklearn.model_selection import train_test_split

data = pd.read_csv('/content/drive/MyDrive/Concept of AI -- week
5/student.csv')
```

1. Print the first 5 and last 5 rows of the dataset

```
print("Top 5 rows of the dataset:")
print(data.head(5))
```

Top 5 rows of the dataset:

	Math	Reading	Writing
0	48	68	63
1	62	81	72
2	79	80	78
3	76	83	79
4	59	64	62

```
print("Bottom 5 rows of the dataset:")
print(data.tail(5))
```

Bottom 5 rows of the dataset:

	Math	Reading	Writing
995	72	74	70
996	73	86	90
997	89	87	94
998	83	82	78
999	66	66	72

2. Print the information of the dataset

```
print("Information about the dataset:")
print(data.info())
```

Information about the dataset:

<class 'pandas.core.frame.DataFrame'>

RangeIndex: 1000 entries, 0 to 999

Data columns (total 3 columns):

#	Column	Non-Null Count	Dtype
0	Math	1000 non-null	int64
1	Reading	1000 non-null	int64
2	Writing	1000 non-null	int64

dtypes: int64(3)

memory usage: 23.6 KB

None

3. Gather descriptive statistics about the dataset

```
print("Descriptive statistics of the dataset:")
print(data.describe())
```

Descriptive statistics of the dataset:

	Math	Reading	Writing
count	1000.000000	1000.000000	1000.000000
mean	67.290000	69.872000	68.616000
std	15.085008	14.657027	15.241287
min	13.000000	19.000000	14.000000
25%	58.000000	60.750000	58.000000
50%	68.000000	70.000000	69.500000
75%	78.000000	81.000000	79.000000
max	100.000000	100.000000	100.000000

```
# 4. Split the dataset into Features (X) and Labels (Y)
# Features: Math and Reading marks, Labels: Writing marks
X = data[['Math', 'Reading']]
Y = data['Writing']
```

```
print("Feature X:")
print(X.head())
```

```
print("Label Y:")
print(Y.head())
```

Feature X:

	Math	Reading
0	48	68
1	62	81
2	79	80
3	76	83
4	59	64

Label Y:

0	63
1	72
2	78
3	79
4	62

Name: Writing, dtype: int64

```
# Step 1: Convert Features (X) and Labels (Y) into matrices
X_matrix = X.to_numpy().T # Convert X to a numpy array and transpose
it to shape (d, n)
Y_matrix = Y.to_numpy() #Convert Y to a numpy array with shape (n, )
```

```
# Step 2: Create a weights matrix (W) initialized to zeros
# Number of features = d (rows of X_matrix)
num_features = X_matrix.shape[0]
W_matrix = np.zeros((num_features,))
```

```
# Print the matrices for verification
print("Feature Matrix (X):\n", X_matrix)
```

```
print("Label Matrix (Y):\n", Y_matrix)
print("Weights Matrix (W):\n", W_matrix)
```

Feature Matrix (X):

```
[[48 62 79 ... 89 83 66]
 [68 81 80 ... 87 82 66]]
```

Label Matrix (Y):

```
[ 63  72  78  79  62  85  83  41  80  77  64  90  45  77  70  46  76
44
 85  72  53  66  75  49  84  83  68  66  77  78  74  83  72  65  46
66
 50  79  68  46  86  70  61  53  72  75  50  77 100  81 100  87  78
48
 50  44  48  43  67  78  58  91  92  78  42  85  73  83  61  58  60
55
 48  62  68  59  62  48  74  63  80  79  73  79  45  67  89  77  81
88
 53  68  79  77  63  73  60  67 100  79  26  51  80  57  41  78  68
49
 76  41  71  77  89  86  55  80  56  74  85  80  73  74  86  56  53
44
 41  59  71  81  74  78  67  53  56  75  82  79  99  76  59  96  75
61
 56  88  65 100  79  55  61  83  74  59  54  47  82  74  59  74  84
59
 43  65  61  78  84  73  73  92  63  72  61  59  70  87  78  65  73
62
 69  55  73  63  67  86  78  85  83  80  60  90  56  70  55  80  82
60
 78  76  94  75  68  71  85  46  58  46  84  58  57  59  77  63  68
99
 48  91  57  80  46  75  59  87  82  79  66  68  66  61  66  63  72
73
 77  84  83  42  72  76  76  39  74  43  63  74  52  31  65  45  87
63
 51  82  86  76  27  70  79  66  61  62  47  17  65  76  75  66  59
61
 93  40  66  43  71  64  55  86  65  70  65  53  49  67  76  95  76
48
 60  53  69  78  62  66  51  52  46  42  77  57 100  84  68  48  72
50
 72  55  72  77  56  94  67  82  75  80  60  73  74  62  53  69  75
60
 58  71  87  74  87  73  78  76  74  55  94  71  76  59  91  57  83
59
 93  64  58  79  96  76  64  70  80  33  95  64  92  34  72  81  57
79
 84  82  54  45  54  62  49  74  59  63  83  62  72  72  65  65  54
78
 82  85  74  83  71  83  77  66  75  52  68  84  67  70  41  91  46
```

[illegible]

```

66 62 99 62 53 57 78 56 87 79 63 87 86 75 70 60 49
41 78 58 75 89 34 60 80 85 73 58 69 74 52 58 79 86 61
68 67 48 65 73 57 73 57 80 85 81 61 69 100 99 92 72 57
44 59 62 93 64 57 72 40 85 60 83 63 74 44 61 74 68 78
50 70 68 82 46 96 100 44 41 95 79 67 52 87 75 61 42 60
57 64 52 68 58 93 75 77 66 63 90 43 65 95 86 31 95 52
63 87 70 59 84 79 77 75 66 69 85 63 50 58 80 47 55 61
87 77 54 66 68 54 69 74 81 72 61 76 63 64 73 62 92 69
70 65 53 74 61 80 85 62 80 83 56 76 52 51 74 57 63 61
87 60 54 89 67 56 70 90 94 78 72]

```

Weights Matrix (W):
[0. 0.]

Step 1: Split the dataset

```
X_train, X_test, Y_train, Y_test = train_test_split(X, Y, test_size =
0.2, random_state = 42)
```

Step 2: Verify the split

```
print("\nTraining Features (X_train):")
print(X_train.head())
```

```
print("\nTraining Labels (Y_train):")
print(Y_train.head())
```

```
print("\nTesting Features (X_test):")
print(X_test.head())
```

```
print("\nTesting Labels (Y_test):")
print(Y_test.head())
```

Training Features (X_train):

	Math	Reading
29	64	82
535	62	70
695	36	21
557	81	70
836	82	86

Training Labels (Y_train):

29	78
----	----

```
535     67
695     25
557     71
836     87
Name: Writing, dtype: int64
```

Testing Features (X_test):

```
      Math  Reading
521     63      69
737     42      37
740     69      62
660     69      59
411    100      92
```

Testing Labels (Y_test):

```
521     69
737     41
740     57
660     56
411     88
```

Name: Writing, dtype: int64

#Define the cost function

```
def cost_function(X, Y, W):
    """ Parameters:
        This function finds the Mean Square Error.
        Input parameters:
        X: Feature Matrix
        Y: Target Matrix
        W: Weight Matrix
        Output Parameters:
        cost: accumulated mean square error.
    """
    n = len(Y)
    Y_pred = np.dot(X, W)
    cost = (1/n) * np.sum((Y_pred - Y) ** 2)
    return cost

def gradient_descent(X, Y, W, alpha, iterations):
    """
        Perform gradient descent to optimize weights.
        Parameters:
        X: Feature matrix (n_samples x n_features)
        Y: Target vector (n_samples)
        W: Weight vector (n_features)
        alpha: Learning rate
        iterations: Number of iterations

        Returns:
        W_update: Optimized weights
    """
```

```

cost_history: History of cost function values
"""
cost_history = []
m = len(Y)

for i in range(iterations):
    # Step 1: Compute predictions
    Y_pred = np.dot(X, W)

    # Step 2: Compute gradient
    gradient = (1/m) * np.dot(X.T, (Y_pred - Y))

    # Step 3: Update weights
    W -= alpha * gradient

    # Step 4: Compute and store cost
    cost = cost_function(X, Y, W)
    cost_history.append(cost)

return W, cost_history

# RMSE Calculation
def rmse(Y, Y_pred):
    """
    Calculate Root Mean Squared Error.
    """
    return np.sqrt(np.mean((Y - Y_pred) ** 2))

# R-Squared Calculation
def r2(Y, Y_pred):
    """
    Calculate R-squared value.
    """

    mean_y = np.mean(Y)
    ss_tot = np.sum((Y - mean_y) ** 2)
    ss_red = np.sum((Y - Y_pred) ** 2)
    return 1 - (ss_red / ss_tot)

def main():
    data = pd.read_csv('/content/drive/MyDrive/Concept of AI -- week
5/student.csv')

    # Step 2: Split the data into features (X) and target (Y)
    X = data[['Math', 'Reading']].values
    Y = data['Writing'].values

    # Step 3: Split the data into training and test sets
    X_train, X_test, Y_train, Y_test = train_test_split(X, Y,
test_size=0.2, random_state = 42)

```

```

# Step 4: Initialize weights (W), learning rate, and iterations
W = np.zeros(X.shape[1])
alpha = 0.00001
iterations = 1000

# Step 5: Perform Gradient Descent
W_optimal, cost_history = gradient_descent(X_train, Y_train, W,
alpha, iterations)

# Step 6: Make predictions on the test set
Y_pred = np.dot(X_test, W_optimal)

# Step 7: Evaluate the model using RMSE and R-Squared
model_rmse = rmse(Y_test, Y_pred)
model_r2 = r2(Y_test, Y_pred)

print("Final Weights:", W_optimal)
print("Cost History (First 10 iterations):", cost_history[:10])
print("RMSE on the Set:", model_rmse )
print("R-squared on the Set:", model_r2)

# Execute the main function
if __name__ == "__main__":
    main()

```

```

Final Weights: [0.34811659 0.64614558]
Cost History (First 10 iterations): [4026.33114156751,
3280.573665199384, 2674.1239989803175, 2180.9589785701155,
1779.9166540166468, 1453.788198601909, 1188.5794521617188,
972.910410590327, 797.5268927198968, 654.9034294649376]
RMSE on the Set: 5.2798239764188635
R-squared on the Set: 0.8886354462786421

```

1. Model Performance: Overfitting, Underfitting, or Acceptable Performance? Key Metrics to Evaluate Performance:

Root Mean Square Error (RMSE):

-- Low RMSE indicates the model predicts the target values well.

R-squared (R^2):

-- A value close to 1 implies the model explains most of the variance in the target variable.

Scenario Evaluation:

- Overfitting:

-- If R^2 on the training set is significantly higher than R^2 on the test set, the model memorizes the training data and fails to generalize.

- Underfitting:

-- Both training and test R^2 are low, indicating the model is too simple to capture the relationship.

- Acceptable Performance:

-- The model achieves a similar R^2 score for both training and test sets, with a low RMSE.

Observation from the Results:

-- If the RMSE on the test set is small (e.g., ~4.57) and the R^2 is high (~0.85), it suggests the model performs well and generalizes effectively. Comparing training and test performance metrics (which can be included in the main function) will confirm whether the model overfits, underfits, or is acceptable.

1. Experimenting with Learning Rate: Impact of Learning Rate (α):

Higher Learning Rate: Convergence may happen faster, but too high a learning rate can cause the model to diverge or oscillate. Lower Learning Rate: The model converges slowly and may require more iterations, but it ensures stability. Experimentation:

Learning Rate (α) = 0.1: Likely to converge faster but might overshoot the minimum or diverge.

Learning Rate (α) = 0.0001: Likely stable but may take many iterations to converge. Default

Learning Rate (α = 0.00001): A balanced learning rate providing stable and accurate convergence.

```
# Set hyperparameters
iterations = 1000 # Define the number of iterations for gradient
descent
```

```
# Experiment with Learning Rates
for alpha in [0.1, 0.01, 0.0001, 0.00001]:
    W_optimal, cost_history = gradient_descent(X_train, Y_train,
np.zeros(X_train.shape[1]), alpha, iterations)
    Y_pred = np.dot(X_test, W_optimal)
    model_rmse = rmse(Y_test, Y_pred)
    model_r2 = r2(Y_test, Y_pred)
    print(f"Learning Rate: {alpha}")
    print(f"RMSE on Test Set: {model_rmse}")
    print(f"R-Squared on Test Set: {model_r2}")
    print("-" * 40)
```

```
<ipython-input-15-e64a3002b6cd>:26: RuntimeWarning: invalid value
encountered in subtract
    W -= alpha * gradient
```

```
Learning Rate: 0.1
RMSE on Test Set: nan
R-Squared on Test Set: 1.0
-----
```

```
/usr/local/lib/python3.10/dist-packages/numpy/core/_methods.py:49:
RuntimeWarning: overflow encountered in reduce
    return umr_sum(a, axis, dtype, out, keepdims, initial, where)
<ipython-input-15-e64a3002b6cd>:26: RuntimeWarning: invalid value
encountered in subtract
    W -= alpha * gradient
```

```
Learning Rate: 0.01
RMSE on Test Set: nan
R-Squared on Test Set: 1.0
```

```
-----
Learning Rate: 0.0001
RMSE on Test Set: 4.792607360540954
R-Squared on Test Set: 0.908240340333986
```

```
-----
Learning Rate: 1e-05
RMSE on Test Set: 5.2798239764188635
R-Squared on Test Set: 0.8886354462786421
-----
```